

lab1-ml

November 10, 2015

# Contents

<b>1</b>	<b>Tarea N°1 - Máquinas de Aprendizaje - ILI393</b>	<b>3</b>
1.0.1	Martín Villanueva A.	3
1.1	Introducción	3
1.2	Parte 1 - Regresión Lineal	3
1.2.1	1a) Gradiente Descendente Batch	3
1.2.2	Raw data	3
1.2.3	Rescaled data	5
1.2.4	Normalized data	7
1.2.5	1b) Gradiente Descendente Online	9
1.2.6	Raw data	9
1.2.7	Rescaled data	11
1.2.8	Normalized data	13
1.2.9	1c) Newton Raphson	15
1.2.10	Raw data	15
1.2.11	Rescaled data	16
1.2.12	Normalized data	16
1.2.13	2a) Gradiente Descendente Batch	17
1.2.14	2b) Gradiente Descendente Online	17
1.2.15	2c) Newton Raphson	17
1.2.16	3) Locally weighted linear regression	17
1.2.17	Raw data	17
1.2.18	Rescaled data	19
1.2.19	Normalized data	21
1.2.20	4)	23
1.2.21	5)	23
1.3	Parte 2 - Regresión Logística	23
1.3.1	1a) Gradiente ascendente online	23
1.3.2	Raw data	23
1.3.3	Rescaled data	26
1.3.4	Normalized data	28
1.3.5	1b) Newton Raphson	30
1.3.6	Raw data	30
1.3.7	Rescaled data	31
1.3.8	Normalized data	31
1.3.9	2a) Gradiente Ascendente Online	32
1.3.10	2b) Newton Raphson	32
1.3.11	3)	32
1.4	Conclusiones	32
1.5	Anexos	32
1.5.1	Configuración del notebook	32
1.5.2	Métricas de error para regresión lineal	32
1.5.3	Implementación de algoritmos de regresión lineal	33

1.5.4	Implementación de locally weighted linear regression . . . . .	34
1.5.5	Métricas de error para regresión logística . . . . .	34
1.5.6	Implementación de algoritmo de regresión logística . . . . .	34
1.5.7	Funciones para manejo de la data . . . . .	35
1.5.8	Funciones para Cross-Validation . . . . .	36
1.5.9	Funciones complementarias (Helpers) para obtener resultados . . . . .	37

# Chapter 1

## Tarea N°1 - Máquinas de Aprendizaje - ILI393

1.0.1 Martín Villanueva A.

### 1.1 Introducción

En es primera tarea, se tiene como objetivo la implementación y testeo de algoritmos para regresión lineal y regresión logística. En ambos casos la búsqueda de los mejores parámetros del modelo se realiza por medio de *Gradiente Descendente* (batch y online) y *Newton-Raphson*. Para la correcta selección de los *hiperparámetros* se realiza 5-fold crossvalidation, intentando de este modo que los modelos resultantes no caigan en problemas de *overfitting*.

### 1.2 Parte 1 - Regresión Lineal


```
In [29]: #alphas to try on raw data
         alphas1 = np.linspace(4.0e-7, 4.5e-7, 5, endpoint=True)
         #alphas to try on rescaled and normalized data
         alphas2 = np.array([1.0e-3, 0.8e-3, 0.6e-3, 0.4e-3, 0.2e-3])
```

#### 1.2.1 1a) Gradiente Descendente Batch

#### 1.2.2 Raw data

```
In [42]: solve_regression(gd_batch, 'linear', params=alphas1, show=[0,14])
```

```
#####
Dataset: 0
Best alpha: 4.5e-07
```



Training error: 291.529613576

Testing error: 224.340517783

N° iterations: 64


Beta: [ 0.00759594 0.25955161 0.02174389 -0.0171124 -0.01772088 0.01810715  
0.13438221 -0.04736688 0.12759068 -0.03385616 0.01151583 0.0052043  
0.0076039 ]

#####

#####

Dataset: 14


Best alpha: 4.5e-07



```

Training error: 236.271311927
Testing error: 424.701874
N° iterations: 53
Beta: [ 0.00555882  0.25899926  0.02076567 -0.00916217 -0.00514845  0.00739162
        0.10491754 -0.03100194  0.1036039   0.02771736  0.01037837  0.00402665
        0.0064281 ]
#####

```




### 1.2.3 Rescaled data

```

In [44]: solve_regression(gd_batch, 'linear', params=alphas2, data_func=rescale, show=[0,14])

#####
Dataset: 0
Best alpha: 0.001

```




Training error: 20.6182739329  
 Testing error: 42.5328851194  
 N° iterations: 530  
 Beta: [ 39.27502848 -13.60536851 17.88310933 -16.95659792 -13.32113539  
 21.00349746 14.26238428 -17.86462666 10.29693403 -11.62685618  
 3.58148223 4.97282493 18.67876439]

#####

#####


Dataset: 14  
 Best alpha: 0.001



```

Training error: 21.9788628555
Testing error: 91.1000283764
N° iterations: 548
Beta: [ 31.86746216 -11.73989936  18.4682303  -15.58587584 -10.42352654
        16.69472148  16.37413566 -14.74326499   7.4469539   -7.23761563
         4.18365522   3.54434565  18.06400087]
#####

```



#### 1.2.4 Normalized data


```

In [45]: solve_regression(gd_batch, 'linear', params=alphas2, data_func=normalize, show=[0,14])

#####
Dataset: 0
Best alpha: 0.001

```





Training error: 13.4194733738

Testing error: 36.5826804086

N° iterations: 137

Beta: [ 42.22626947 -4.02131789 4.23055468 -3.50410785 -3.8804333

6.72285551 3.1979057 -4.13343304 -0.61328431 -3.02687997


1.98386787 1.43159348 3.59464352]

#####

#####

Dataset: 14


Best alpha: 0.001



```

Training error: 16.3098751254
Testing error: 65.7517896491
N° iterations: 96
Beta: [ 40.87319656 -3.87689122  4.66900552 -3.56542152 -3.20285147
        5.02062112  2.99775041 -3.08497395  0.46571746 -1.71364208
        1.8960333  1.21081326  3.61273617]
#####

```



## 1.2.5 1b) Gradiente Descendente Online


### 1.2.6 Raw data

```
In [46]: solve_regression(gd_online, 'linear', params=alphas1, show=[0,14])
```

```

#####
Dataset: 0
Best alpha: 4.5e-07

```



Training error: 292.641332227

Testing error: 216.286212008

N° iterations: 66


Beta: [ 0.00777387 0.26225985 0.0228718 -0.01769797 -0.01559905 0.0187834  
0.13911965 -0.05018793 0.13196834 -0.03041284 0.01187493 0.00535709  
0.00780838]

#####

#####

Dataset: 14


Best alpha: 4.5e-07



```

Training error: 236.224475078
Testing error: 425.608168754
N° iterations: 54
Beta: [ 0.0056634  0.25890305  0.02099604 -0.00953796 -0.00318058  0.00795135
 0.10658788 -0.03180157  0.10336748  0.02722773  0.01067373  0.00410388
 0.00649339]
#####

```



### 1.2.7 Rescaled data

```


In [47]: solve_regression(gd_online, 'linear', params=alphas2, data_func=rescale, show=[0,14])

```

```


#####
Dataset: 0
Best alpha: 0.001

```



Training error: 20.6132205513  
 Testing error: 42.5978431746  
 N° iterations: 531  
 Beta: [ 39.25933591 -13.60674956 17.89442223 -16.95968836 -13.34024487  
 21.02057345 14.28707965 -17.853764 10.30068558 -11.61868158  
 3.58670305 4.99221433 18.68272861]  
 #####


#####  
 Dataset: 14  
 Best alpha: 0.001



```

Training error: 21.9786834767
Testing error: 90.8722318343
N° iterations: 548
Beta: [ 31.87057958 -11.74525815  18.46781379 -15.57800553 -10.41308722
        16.7220343  16.37303632 -14.73147343   7.43757436  -7.23999133
         4.19564227   3.56494144  18.04271745]
#####

```




### 1.2.8 Normalized data

```

In [48]: solve_regression(gd_online, 'linear', params=alphas2, data_func=normalize, show=[0,14])

#####
Dataset: 0
Best alpha: 0.001

```



Training error: 13.4160931261

Testing error: 36.4127700422

N° iterations: 137

Beta: [ 42.22489559 -4.01812148 4.22806184 -3.50756328 -3.86158705

6.74149223 3.19810492 -4.11134324 -0.61343112 -3.0115829


1.97701226 1.44230261 3.57298429]

#####

#####

Dataset: 14


Best alpha: 0.001



```

Training error: 16.3153752729
Testing error: 65.4658388653
N° iterations: 95
Beta: [ 40.88418912 -3.86109769  4.6767013  -3.52334001 -3.23350274
        5.05247161  3.00863291 -3.07361825  0.48458418 -1.72351009
        1.90061769  1.19267068  3.59741601]
#####

```



### 1.2.9 1c) Newton Raphson

#### 1.2.10 Raw data

```
In [33]: solve_regression(nr_linear, 'linear', show=[0,14])
```

```

#####
Dataset: 0
Training error: 12.690235978

```



```

Testing error: 29.8692128464
N° iterations: 2
Beta: [ 27.77312374 -0.19196441  3.64271996 -2.61978866 -0.04940511
        3.27558724  0.8950743  -0.94639251 -0.03448292 -0.13677948
        2.89862662 12.4397042  13.71043458]

```

```
#####
```

```
#####
```

```

Dataset: 14
Training error: 15.5168045543
Testing error: 22.256882111
N° iterations: 2
Beta: [ 2.92044009e+01 -2.45567403e-01  3.72846884e+00 -2.50624227e+00
        -3.72800837e-02  2.92327199e+00  8.33168001e-01 -6.98788024e-01
        -1.71242668e-02 -1.04831696e-01  2.76379484e+00  1.24136911e+01
        1.31720844e+01]

```

```
#####
```

### 1.2.11 Rescaled data

```
In [35]: solve_regression(nr_linear, 'linear', data_func=rescale, show=[0,14])
```

```
#####
```

```

Dataset: 0
Training error: 12.690235978
Testing error: 70.8230012379
N° iterations: 2
Beta: [ 37.68776506 -21.30804932  18.21359978 -15.71873198 -15.80963497
        45.85822131  16.11133741 -17.03506525 -10.86211949 -13.67794787
        5.79725325 12.4397042  24.27432442]

```

```
#####
```

```
#####
```

```

Dataset: 14
Training error: 15.5168045543
Testing error: 153.683962628
N° iterations: 2
Beta: [ 32.95912586 -27.25798171  18.64234419 -15.03745362 -11.92962679
        29.23271995  18.32969603 -12.57818444 -5.37701976 -10.48316963
        5.52758969 12.41369113  23.32117544]

```

```
#####
```

### 1.2.12 Normalized data

```
In [36]: solve_regression(nr_linear, 'linear', data_func=normalize, show=[0,14])
```

```
#####
```

```

Dataset: 0
Training error: 12.690235978
Testing error: 39.1167896725
N° iterations: 2
Beta: [ 42.22724  -3.99216233  4.26196505 -3.44486437 -3.93375017
        8.28671753  3.23536861 -4.20056151 -2.37742955 -3.2700182

```

```

2.44555845  1.89445678  3.72993749]
#####

#####
Dataset: 14
Training error: 15.5168045543
Testing error: 61.1820073443
N° iterations: 2
Beta: [ 40.89617333 -4.54027498  4.63264649 -3.14293395 -2.94434592
  6.47945038  3.19280105 -3.05719739 -1.09358413 -2.01526648
  2.3439993  1.82778689  3.80201253]
#####

```

### 1.2.13 2a) Gradiente Descendente Batch

### 1.2.14 2b) Gradiente Descendente Online

### 1.2.15 2c) Newton Raphson

### 1.2.16 3) Locally weighted linear regression

```

In [49]: taus1 = np.linspace(30.,200.,5, endpoint=True)
        taus2 = np.linspace(1.,10.,5, endpoint=True)

```

### 1.2.17 Raw data

```


In [57]: solve_weighted(taus1, show=[0,14])

```

```


#####
Dataset: 0
Best tau: 30.0

```

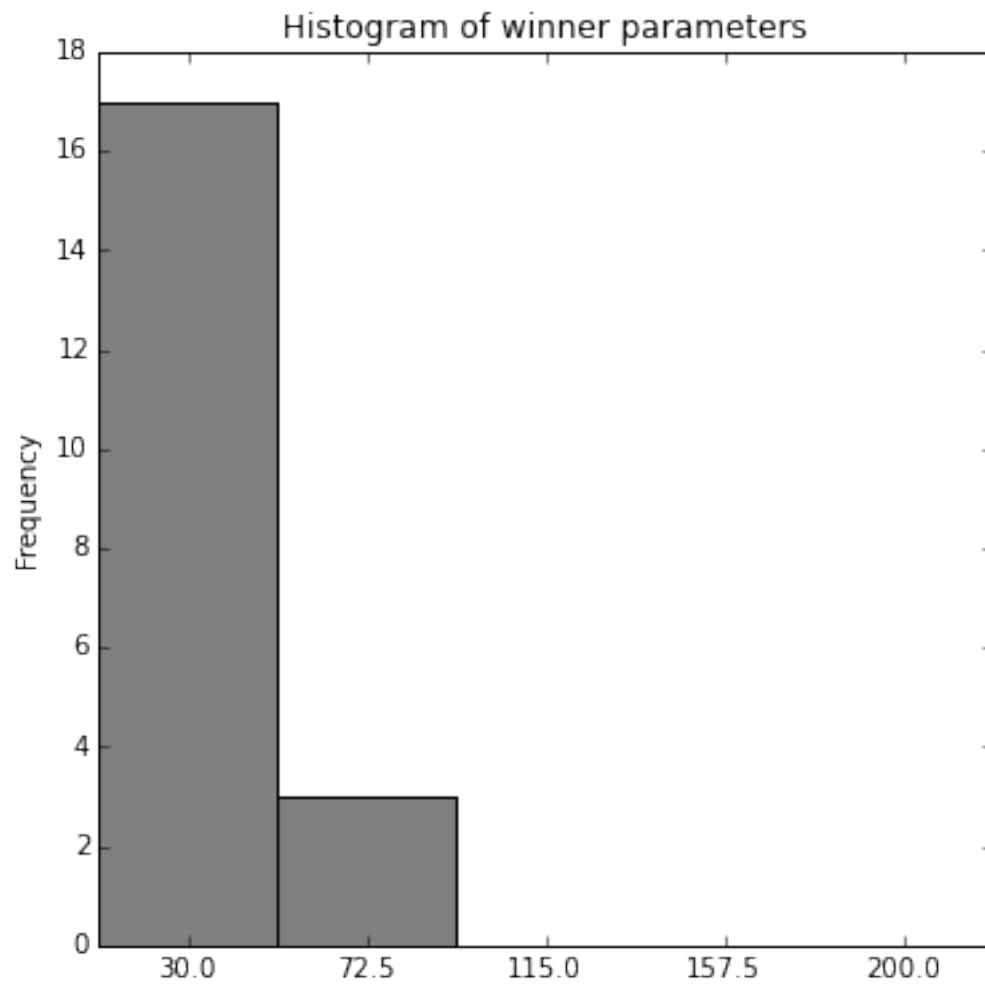


Training error (weighted): 0.185160073025  
 Testing error (weighted): 9.27828692439  
 Training error: 240.503072277  
 Testing error: 332.258498635  
 #####

#####  
 Dataset: 14  
 Best tau: 30.0




Training error (weighted): 0.26539563072  
 Testing error (weighted): 2.51726674875  
 Training error: 99.4681361427  
 Testing error: 136.928680467  
 #####



### 1.2.18 Rescaled data

```
In [58]: solve_weighted(taus2, data_func=rescale, show=[0,14])
```

```
#####  
Dataset: 0  
Best tau: 1.0
```



Training error (weighted): 5.8374635318

Testing error (weighted): 34.5095818121

Training error: 12.9492812413


Testing error: 70.2835184179

#####

#####

Dataset: 14


Best tau: 1.0



```

Training error (weighted): 6.62249985344
Testing error (weighted): 66.2273010366
Training error: 16.0705146463
Testing error: 159.675444498
#####

```




### 1.2.19 Normalized data

```

In [59]: solve_weighted(taus2, data_func=rescale, show=[0,14])

#####
Dataset: 0
Best tau: 1.0

```



Training error (weighted): 5.8374635318

Testing error (weighted): 34.5095818121

Training error: 12.9492812413


Testing error: 70.2835184179

#####

#####

Dataset: 14

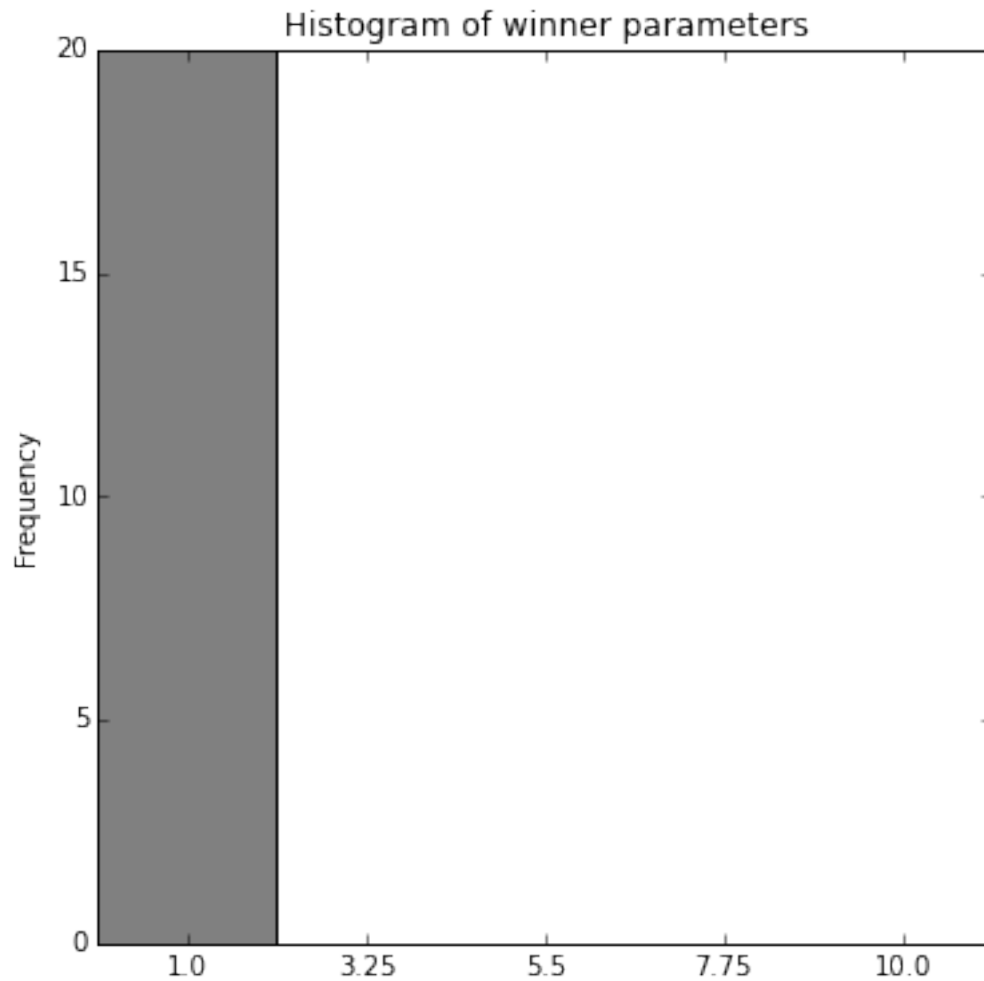
Best tau: 1.0



```

Training error (weighted): 6.62249985344
Testing error (weighted): 66.2273010366
Training error: 16.0705146463
Testing error: 159.675444498
#####

```



1.2.20 4)

1.2.21 5)

## 1.3 Parte 2 - Regresión Logística

### 1.3.1 1a) Gradiente ascendente online

```

In [61]: #alphas to try on ascent gradient stochastic
         alphas3 = np.linspace(1e-2, 1e-3, 5, endpoint=True)

```

### 1.3.2 Raw data

```

In [73]: solve_regression(gd_stochastic, 'logistic', params=alphas3, show=[0,14])


```



#####

Dataset: 0

Best alpha: 0.00775



Training error: 0.0444444444444

Testing error: 0.0666666666667

N° iterations: 13625


Beta: [ -6.07258811e+00 7.03329870e+00 -8.26724873e+00 -2.25430210e-01  
1.24680827e+02 -2.80636719e+02 3.07024050e+00]

#####

#####

Dataset: 14

Best alpha: 0.0055




Training error: 0.155555555556

Testing error: 0.166666666667

N° iterations: 100000

Beta: [ 85.582449      5.51768624    -6.92887744      0.90383805    137.34901635  
 -394.48396529    0.9807746 ]

#####




### 1.3.3 Rescaled data

```
In [71]: solve_regression(gd_stochastic, 'logistic', params=alphas3, data_func=rescale, show=[0,14])
```

```
#####
```

```
Dataset: 0
```

```
Best alpha: 0.01
```



Training error: 0.0666666666667

Testing error: 0.1

N° iterations: 364


Beta: [-0.89067992 1.89240682 -3.06533656 -0.49413702 8.38805164 -3.67234941  
1.39561496]

#####

#####

Dataset: 14


Best alpha: 0.01



```

Training error: 0.0888888888889
Testing error: 0.1
N° iterations: 341
Beta: [-0.86739457  2.20861945 -2.52849556 -0.58885793  8.37423558 -3.46780645
        0.54455779]
#####

```




### 1.3.4 Normalized data

```

In [72]: solve_regression(gd_stochastic, 'logistic', params=alphas3, data_func=normalize, show=[0,14])

#####
Dataset: 0
Best alpha: 0.00325

```



Training error: 0.0222222222222

Testing error: 0.0333333333333

N° iterations: 345


Beta: [-0.2385682 0.81850858 -1.27981156 -0.11349321 3.88872458 -2.76731914  
0.40282272]

#####

#####

Dataset: 14


Best alpha: 0.0055



```

Training error: 0.0444444444444
Testing error: 0.0666666666667
N° iterations: 290
Beta: [ 0.20829953  1.02332025 -1.47941872 -0.31923802  4.65487619 -3.11664099
 0.09574176]
#####

```



### 1.3.5 1b) Newton Raphson

### 1.3.6 Raw data

```
In [67]: solve_regression(nr_logistic, 'logistic', data_func=normalize, show=[0,14])
```

```

#####
Dataset: 0
Training error: 0.0
Testing error: 0.0333333333333

```

```

N° iterations: 17
Beta: [ -10.17632429  51.90495222 -223.06462833 -72.5822921  446.92953071
        -290.39247559   5.29862312]

```

```
#####
```

```
#####
```

```

Dataset: 14
Training error: 0.022222222222222
Testing error: 0.033333333333333
N° iterations: 14
Beta: [  6.21575176  18.14870684 -124.70691036 -40.65161461  235.09789749
        -145.40751338  -3.76500934]

```

```
#####
```

### 1.3.7 Rescaled data

```
In [69]: solve_regression(nr_logistic, 'logistic', data_func=rescale, show=[0,14])
```

```
#####
```

```

Dataset: 0
Training error: 0.0
Testing error: 0.1
N° iterations: 17
Beta: [  73.98701993  180.88595375 -814.87368456 -183.34781483  1865.9231994
        -602.60602669   22.16612173]

```

```
#####
```

```
#####
```

```

Dataset: 14
Training error: 0.022222222222222
Testing error: 0.133333333333333
N° iterations: 14
Beta: [  66.60527006  62.6697612 -440.52106091 -98.52834405  948.8637574
        -296.81184378 -17.07077338]

```

```
#####
```

### 1.3.8 Normalized data

```
In [70]: solve_regression(nr_logistic, 'logistic', data_func=normalize, show=[0,14])
```

```
#####
```

```

Dataset: 0
Training error: 0.0
Testing error: 0.033333333333333
N° iterations: 17
Beta: [ -10.17632429  51.90495222 -223.06462833 -72.5822921  446.92953071
        -290.39247559   5.29862312]

```

```
#####
```

```
#####
```

```

Dataset: 14
Training error: 0.022222222222222

```



```

Testing error: 0.0333333333333
N° iterations: 14
Beta: [ 6.21575176 18.14870684 -124.70691036 -40.65161461 235.09789749
       -145.40751338 -3.76500934]
#####

```

### 1.3.9 2a) Gradiente Ascendente Online

### 1.3.10 2b) Newton Raphson

### 1.3.11 3)

## 1.4 Conclusiones

## 1.5 Anexos

En la siguiente sección se encuentra todo el código necesario para reproducir cada uno de los resultados mostrados anteriormente. Para poder ejecutar el código en el informe, se debe en primer lugar ejecutar las celdas de código presentes en este anexo.

### 1.5.1 Configuración del notebook

```

In [17]: #notebook settings
         %matplotlib inline

         #import some useful libraries and utilities
         import numpy as np
         import matplotlib.pyplot as plt
         from sklearn.cross_validation import KFold
         from sklearn.cross_validation import KFold

         #avoid numpy warning (they are handled correctly)
         np.seterr('ignore')

         #setting some paths
         path1 = './cereales/'
         #data directory
         path2 = './credit/'

```

### 1.5.2 Métricas de error para regresión lineal

```

In [18]: #overall cost function for linear regression
         def J(X, y, beta):
             f = np.dot(X,beta)
             diff = f-y
             return 0.5*np.dot(diff,diff)

         #mean squared error for linear regression
         def mse(X, y, beta):
             M,_ = X.shape
             f = np.dot(X,beta)
             diff = f-y
             return (1./(M-1))*np.dot(diff,diff)

```

### 1.5.3 Implementación de algoritmos de regresión lineal

```
In [32]: #batch gradient descent for linear regression
def gd_batch(X, y, alpha, eps=1e-3, max_iter=100000):
    M,N = X.shape
    beta = np.zeros(N)
    J1 = J(X,y,beta)
    for i in xrange(max_iter):
        J0 = J1
        f = np.dot(X,beta)
        dJ = np.dot(X.T,f-y)
        beta -= alpha*dJ
        J1 = J(X,y,beta)
        if np.abs(J1-J0)/J0 < eps:
            break
    return (beta,i+1)

#online gradient descent for linear regression
def gd_online(X, y, alpha, eps=1e-3, max_iter=100000):
    M,N = X.shape
    beta = np.zeros(N)
    J1 = J(X,y,beta)
    for i in xrange(max_iter):
        J0 = J1
        for m in xrange(M):
            beta -= alpha*(np.dot(X[m],beta)-y[m])*X[m]
        J1 = J(X,y,beta)
        if np.abs(J1-J0)/J0 < eps: break
    return (beta,i+1)

#Newton-Raphson method for linear regression
def nr_linear(X, y, eps=1e-5, max_iter=100000):
    M,N = X.shape
    beta = np.zeros(N)
    J1 = J(X,y,beta)
    Hess = np.dot(X.T,X)
    for i in xrange(max_iter):
        J0 = J1
        f = np.dot(X,beta)
        dJ = np.dot(X.T,f-y)
        beta -= np.linalg.solve(Hess, dJ)
        J1 = J(X,y,beta)
        if np.abs(J1-J0)/J0 < eps: break
    return (beta,i+1)
```

**\*\* Comentarios de implementación:** \* Para todos los algoritmos existen básicamente dos criterios de salida. El primero es cuando el error relativo es menor a *eps*, vale decir, cuando la función de error esta cambiando muy poco de iteración en iteración. El segundo es el número máximo de iteraciones, mas que nada para detener algoritmos que no pueden cumplir con el criterio del error (learning rates muy altos por ejemplo). \* Todos los *starting guess* son el vector zeros. Esto para reproducir y comparar resultados de manera adecuada. \* En vez de invertir la matriz Hessiana en Newton-Raphson, se opta por resolver el sistema lineal asociado, por razones de estabilidad numérica.

### 1.5.4 Implementación de locally weighted linear regression

In [53]: *#compute weights for all samples in X matrix, respect to x0*

```
def weight(X, x0, tau):
    Diff = X - x0
    Diff *= Diff
    return np.exp(-1*np.sum(Diff,axis=1)/(2.*tau**2))
```

*#weighted cost function*

```
def wJ(X, y, beta, w):
    f = np.dot(X,beta)
    diff = f-y
    diff **=2
    return 0.5*np.dot(w,diff)
```

*#weighted mean squared error*

```
def wmse(X, y, beta, w):
    M,_ = X.shape
    f = np.dot(X,beta)
    diff = f-y
    diff **=2
    return (1./(M-1))*np.dot(w,diff)
```

*#find best beta for locally weighted linear regression*

```
def min_weighted(X, y, w):
    W = np.diag(w)
    M = np.dot(X.T, np.dot(W, X))
    b = np.dot(X.T, np.dot(W, y))
    return np.linalg.solve(M,b)
```

### 1.5.5 Métricas de error para regresión logística

In [20]: *#log likelihood function for logistic regression*

```
"""
Computing l this way, make it more stable numerically (no overflows en exp)
"""
```

```
def l(X, y, beta):
    y1_mask = y.astype(bool)
    y0_mask = np.logical_not(y1_mask)
    f = sigmoid(np.dot(X,beta))
    return (np.log(f[y1_mask])).sum() + (np.log(1-f[y0_mask])).sum()
```

*#error rate for logistic regression*

```
def error_rate(X, y, beta):
    h = np.round(sigmoid(np.dot(X,beta)))
    h = h.astype(int)
    y = y.astype(int)
    m, = h.shape
    return np.logical_xor(h,y).sum()/np.float(m)
```

### 1.5.6 Implementación de algoritmo de regresión logística

In [66]: *#sigmoid function*

```
def sigmoid(z):
```

```

    return 1./(1.+np.exp(-z))

#stochastic gradient ascent for logistic regression
def gd_stochastic(X, y, alpha, eps=1e-3, max_iter=100000):
    M,N = X.shape
    beta = np.zeros(N)
    l1 = l(X, y, beta)+1.
    for i in xrange(max_iter):
        l0 = l1
        for m in xrange(M):
            beta += alpha*(y[m]-sigmoid(np.dot(X[m],beta)))*X[m]
        l1 = l(X,y,beta)+1.
        if np.abs(l1-l0)/np.abs(l0) < eps: break
    return (beta,i+1)

#Newton-Raphson method for logistic regression
def nr_logistic(X, y, eps=1e-3, max_iter=100000):
    M,N = X.shape
    beta = np.zeros(N)
    l1 = l(X, y, beta)+1.
    for i in xrange(max_iter):
        l0 = l1
        f = sigmoid(np.dot(X,beta))
        W = np.diag(f*(1-f))
        Hess = -1*np.dot(X.T, np.dot(W, X))
        D1 = np.dot(X.T, y-f)
        #when it converges, Hess became singular
        try:
            beta -= np.linalg.solve(Hess, D1)
        except np.linalg.LinAlgError:
            break
        l1 = l(X, y, beta)+1.
        if np.abs(l1-l0)/np.abs(l0) < eps: break
    return (beta,i+1)

```

**Comentarios de implementación:** \* Para ambos algoritmos hay dos criterios de salida. El primero es cuando la función *log verosimilitud* cambia relativamente menor a *eps* en cada iteración (pues es la función que se quiere maximizar). Se tiene en cuenta además que en el óptimo esta función debe ser 0 (en el óptimo la función de *verosimilitud* es 1, pues maximiza la probabilidad para cada dato), por lo que se le suma un 1 para evitar problemas al computar el criterio de salida. El segundo criterio el número máximo de iteraciones. \* Existe un tercer criterio de salida en el método de *Newton-Raphson*. A medida que converge, el vector  $f$  con las probabilidades de pertenecer a la clase 1 de todos los datos, tiene sólo valores cercanos a 0 y 1. Luego al computar la matriz  $W$ , esta empezará a tener filas completas de 0 o valores muy cercanos a 0, y por lo tanto la matriz Hessiana también, y al converger esta matriz se vuelve singular. Para eso se ocupa el manejo de la excepción en caso de existir singularidad.

### 1.5.7 Funciones para manejo de la data

```

In [22]: #Rescale features of M to [a,b] range
def rescale(M, a=0., b=1.):
    #max and min vectors
    maxv = np.max(M, axis=0)
    minv = np.min(M, axis=0)
    return (b-a)*M/(maxv-minv) + (a*maxv-b*minv)/(maxv-minv)

```

```

#Normalize features of M
def normalize(M):
    #mean and standard deviation vectors
    meanv = np.mean(M, axis=0)
    stdv = np.std(M, axis=0)
    return (M-meanv)/stdv

```

### 1.5.8 Funciones para Cross-Validation

In [23]: *""" find the best learning parameter for algorithm, between parameters in params using 5-fold cross validation """*

```

def cross_alpha(X, y, algorithm, error_func, params):
    #creating kfold
    m,n = X.shape
    kf = KFold(m, n_folds=5)
    cv_err = np.empty((5,5))
    i = 0 #index of fold

    for tr_index,ts_index in kf:
        j = 0 #index of parameter
        X_tr, X_ts = X[tr_index], X[ts_index]
        y_tr, y_ts = y[tr_index], y[ts_index]
        for param in params:
            beta,_ = algorithm(X_tr, y_tr, alpha=param)
            cv_err[i,j] = error_func(X_ts, y_ts, beta)
            j += 1
        i += 1

    #arrays with mean cv-error for each alpha
    cv_mean = np.mean(cv_err, axis=0)
    return params[np.argmin(cv_mean)], cv_err

""" find the best band width parameter for locally weighted linear regression, between parameters in params using 5-fold cross validation """
def cross_tau(X, y, params):
    #creating kfolds
    m,n = X.shape
    kf = KFold(m, n_folds=5)
    cv_err = np.zeros((5,5))
    i = 0 #index of fold

    for tr_index,ts_index in kf:
        X_tr, X_ts = X[tr_index], X[ts_index]
        y_tr, y_ts = y[tr_index], y[ts_index]
        j = 0 #index of parameter
        for tau in params:
            for x0 in X_ts:
                w1 = weight(X_tr, x0, tau)
                w2 = weight(X_ts, x0, tau)
                beta = min_weighted(X_tr, y_tr, w1)
                cv_err[i,j] += wmse(X_ts, y_ts, beta, w2)
            cv_err[i,j] /= X_ts.shape[0]
            j +=1

```

```

        i +=1

#arrays with mean costs for each alpha
cv_mean = np.mean(cv_err, axis=0)
return params[np.argmin(cv_mean)], cv_err

```

### 1.5.9 Funciones complementarias (Helpers) para obtener resultados

```

In [56]: """
Function to generate histogram of winners
"""
def make_hist(winners,params):
    winners = np.array(winners)
    freqs = np.zeros(5)
    for i in xrange(5):
        freqs[i] = np.sum(params[i]==winners)

    labels = map(str,params)
    pos = np.arange(len(labels))
    width = 1.0
    fig = plt.figure()
    fig.set_figheight(6)
    fig.set_figwidth(6)
    ax = plt.axes()
    ax.set_xticks(pos + (width / 2))
    ax.set_xticklabels(labels)
    plt.ylabel('Frequency')
    plt.title('Histogram of winner parameters')
    plt.bar(pos, freqs, width, color='0.5')
    plt.show()

"""
Generate solutions for regression problems
(linear and logistic)
"""
def solve_regression(algorithm, kind, params=None, data_func=None, show=None):
    if params is not None:
        winners = list()

    if kind=='linear':
        path = path1+'cereales'
        error_func = mse
    elif kind=='logistic':
        path = path2+'credit'
        error_func = error_rate
    else:
        print "Unknown kind!"
        return -1

    for i in xrange(20):
        #Loading dataset
        tr_file = path+'-tr-{0}.npy'.format(i)
        ts_file = path+'-ts-{0}.npy'.format(i)
        tr_data = np.load(tr_file)

```

```

ts_data = np.load(ts_file)

if data_func is not None:
    X_tr = data_func(tr_data[:, :-1])
else:
    X_tr = tr_data[:, :-1]
y_tr = np.ascontiguousarray(tr_data[:, -1])
#Adding column of 1's
m,n = X_tr.shape
X_tr = np.concatenate((np.ones((m,1)),X_tr),axis=1)

if data_func is not None:
    X_ts = data_func(ts_data[:, :-1])
else:
    X_ts = ts_data[:, :-1]
y_ts = np.ascontiguousarray(ts_data[:, -1])
#Adding column of 1's
m,n = X_ts.shape
X_ts = np.concatenate((np.ones((m,1)),X_ts),axis=1)

if params is not None:
    alpha,cv_err = cross_alpha(X_tr, y_tr, algorithm, error_func, params)
    winners.append(alpha)
    beta,it = algorithm(X_tr, y_tr, alpha)
else:
    beta,it = algorithm(X_tr, y_tr)

if (show is not None) and (i not in show): continue
print "#####\"
print \"Dataset: {0}\".format(i)
if params is not None:
    print 'Best alpha: {0}'.format(alpha)
    fig = plt.figure()
    fig.set_figheight(5)
    fig.set_figwidth(10)
    plt.xlabel('Parameters')
    plt.ylabel('Cross-Validation Errors')
    plt.boxplot(cv_err, showmeans=True, meanline=True)
    plt.xticks([1, 2, 3, 4, 5], map(str,params))
    plt.show()
print 'Training error: {0}'.format(error_func(X_tr,y_tr,beta))
print 'Testing error: {0}'.format(error_func(X_ts,y_ts,beta))
print 'N° iterations: {0}'.format(it)
print 'Beta: {0}'.format(beta)
print "#####\"
print '\n'

if params is not None:
    make_hist(winners,params)

"""
Generate solutions for locally weighted linear regression problems
"""
def solve_weighted(params, data_func=None, show=None):

```

```

#list with winners-alphas
winners = list()

for i in xrange(20):
    #Loading dataset
    tr_file = path1+'cereales-tr-{0}.npy'.format(i)
    ts_file = path1+'cereales-ts-{0}.npy'.format(i)
    tr_data = np.load(tr_file)
    ts_data = np.load(ts_file)

    if data_func is not None:
        X_tr = data_func(tr_data[:, :-1])
    else:
        X_tr = tr_data[:, :-1]
    y_tr = np.ascontiguousarray(tr_data[:, -1])
    #Adding column of 1's
    m,n = X_tr.shape
    X_tr = np.concatenate((np.ones((m,1)),X_tr),axis=1)

    if data_func is not None:
        X_ts = data_func(ts_data[:, :-1])
    else:
        X_ts = ts_data[:, :-1]
    y_ts = np.ascontiguousarray(ts_data[:, -1])
    #Adding column of 1's
    m,n = X_ts.shape
    X_ts = np.concatenate((np.ones((m,1)),X_ts),axis=1)

    tau,cv_err = cross_tau(X_tr, y_tr, params)
    winners.append(tau)
    wtr_err = 0
    wts_err = 0
    tr_err = 0
    ts_err = 0
    for x0 in X_ts:
        w1 = weight(X_tr, x0, tau)
        w2 = weight(X_ts, x0, tau)
        beta = min_weighted(X_tr, y_tr, w1)
        wtr_err += wmse(X_tr, y_tr, beta, w1)
        wts_err += wmse(X_ts, y_ts, beta, w2)
        tr_err += mse(X_tr, y_tr, beta)
        ts_err += mse(X_ts, y_ts, beta)
    M = X_ts.shape[0]
    wtr_err /= M
    wts_err /= M
    tr_err /= M
    ts_err /= M

    if (show is not None) and (i not in show): continue
    print "#####\n"
    print "Dataset: {0}".format(i)
    print 'Best tau: {0}'.format(tau)
    fig = plt.figure()

```



```

fig.set_figheight(5)
fig.set_figwidth(10)
plt.xlabel('Parameters')
plt.ylabel('Cross-Validation Errors')
plt.boxplot(cv_err, showmeans=True, meanline=True)
plt.xticks([1, 2, 3, 4, 5], map(str, params))
plt.show()
print 'Training error (weighted): {0}'.format(wtr_err)
print 'Testing error (weighted): {0}'.format(wts_err)
print 'Training error: {0}'.format(tr_err)
print 'Testing error: {0}'.format(ts_err)
print "#####"
print '\n'
make_hist(winners, params)

```